


11 Apr 25

Simplest "TD" method



move a small step α in the direction of error

$$V(s_t) \leftarrow V(s_t) + \alpha \left[r_{t+1} + \gamma V(s_{t+1}) - \underbrace{V(s_t)}_{\text{estimate of } V(s_t) \text{ but at time } t+1} \right]$$

$V(s_t)$ at time t my prediction as I land at s_t

TD error δ

- TD like MC - do not require complete env only experience (sampling)
- can be fully incremental (bootstrapping)
- can learn even w/o the final outcome.

→ TD prediction

- $\pi \xrightarrow{\text{get}} V_\pi$ or q_π
- no knowledge of p and r but access to real system / sample model.

TD(0)

$$V_{k+1}(s_t) \leftarrow V_k(s_t) + \alpha \left[r_{t+1} + \gamma V_k(s_{t+1}) - V_k(s_t) \right]$$

$V_{k+1}(s_t)$ take k steps of updates

get sample acc to policy π $V_k \rightarrow$ best estimate of V_π


$$V_\pi = E_\pi \left\{ r_{t+1} + \gamma V_\pi(s_{t+1}) \right\}$$

Transition acc to MDP, act acc to best policy.

MC | Exhaustive Search: Go all the way to the end.

TD @ DP : Stop after a step (bootstrapping)

eg. TD Update Example



$$r \ A \rightarrow B : 0$$

$$\alpha : 0.2 \quad \text{Step Size}$$

$$\gamma = 0.9 \quad \text{Discount factor}$$

$$\begin{aligned} V(A) &= V(A) + \alpha [R + \gamma V_B - V_A] \\ &= 0.5 + 0.2 [0.9(0.8) - 0.5] \\ &= 0.544 \end{aligned}$$

→ MC vs TD Updates

Same trajectory, value function by MC vs TD

Batch TD: Take a batch / set of episodes
loop TD
till convergence

MC: first visit

fixed policy.

Consider: A - B

AOBO (terminate)
B1
B0
A1 B0
B1
B0

MC

$$V(A) = \frac{0+1}{2} = 0.5$$

look at AOBO
A1 B0

$$V(B) = \frac{0+1+0+0+1+0}{6} = \frac{2}{6} = \frac{1}{3} = 0.33$$

$$\gamma = 1$$

TD


$$V(A) = E[R + \gamma V(B)] = 1 + \gamma V(B) = 1 + \gamma \left(\frac{1}{3}\right)$$

every time
I go to B

$$= 0.833$$

$$V(B) = E[R] = \frac{1}{3}$$

"implicitly" forming an MDP even though I am only given samples of an MDP




which is correct? happen because finite data.

lot of data \rightarrow both converge to same.

\rightarrow MC - converges to min least squares estimate of return.

TD - certainty equivalence estimate.

\rightarrow SARSA



but Q learning we don't use next action


look at TD evaluation

now

TD control

find the optimal policy.

GPI:



Policy Evaluation : use TD(0)

Policy Improvement : make greedy wrt current value function

Note:

We estimate action values rather than state values in the absence of model.

ϵ -Greedy Policies:

$$a^* \leftarrow \operatorname{argmax}_a Q(s, a)$$

$$\forall a \in A(s):$$

$$\pi(a|s) = \begin{cases} 1 - \epsilon + \frac{\epsilon}{|A(s)|} & \text{if } a = a^* \\ \frac{\epsilon}{|A(s)|} & \text{if } a \neq a^* \end{cases}$$

→ any ϵ greedy policy w.r.t Q following π is an impr over any ϵ -soft policy is assured by the policy improvement theorem.
 $\pi(a|s)$ is at least ϵ for every a

(1) SARSA: On-Policy TD control

samples acc
to policy we
are trying to
evaluate & improve

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t) \right]$$

TD(0) for Q

$$Q(s_{t+1}, a_{t+1}) = 0 \text{ if } s_{t+1} \text{ is terminal}$$

Sarsa Algorithm:

initialize $Q(s, a)$ arbitrarily

For each episode

π : derived from Q

$$a \sim \pi(\cdot) \text{ (eq. } \epsilon \text{ greedy)}$$

For each step in episode

$$s \xrightarrow[r]{a} s'$$

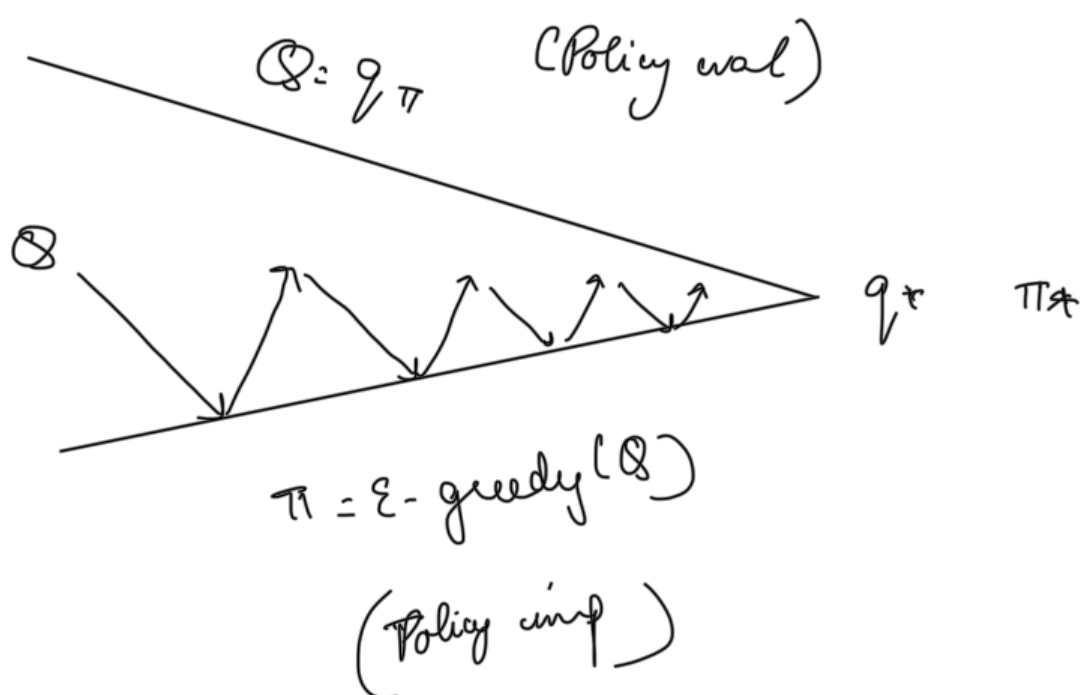
(improvement)
(evaluation)

$$a' = \pi(s')$$

$$Q(s,a) \leftarrow Q(s,a) + \alpha [r + \gamma Q(s',a') - Q(s,a)]$$

$$s \leftarrow s' \quad a \leftarrow a'$$

until s is terminal



Convergence

- all (s,a) vis ∞

- converges in the limit to the greedy policy $\epsilon \rightarrow 0$

(GLIE)

Q-learning

One-step Q-learning

$$\hat{Q}(s_t, a_t) \leftarrow \hat{Q}(s_t, a_t) + \alpha \left[r_{t+1} + \underbrace{\gamma \max_a \hat{Q}(s_{t+1}, a)}_{\text{in SARSA } Q(s_{t+1}, a_{t+1})} - \hat{Q}(s_t, a_t) \right]$$

Temporal Difference

Bellman Optimality Equation:

$$Q^*(s, a) = E \left\{ r_{t+1} + \gamma \max_{a'} Q^*(s_{t+1}, a') \mid s_t = s, a_t = a \right\}$$

Stochastic Averaging Rule:

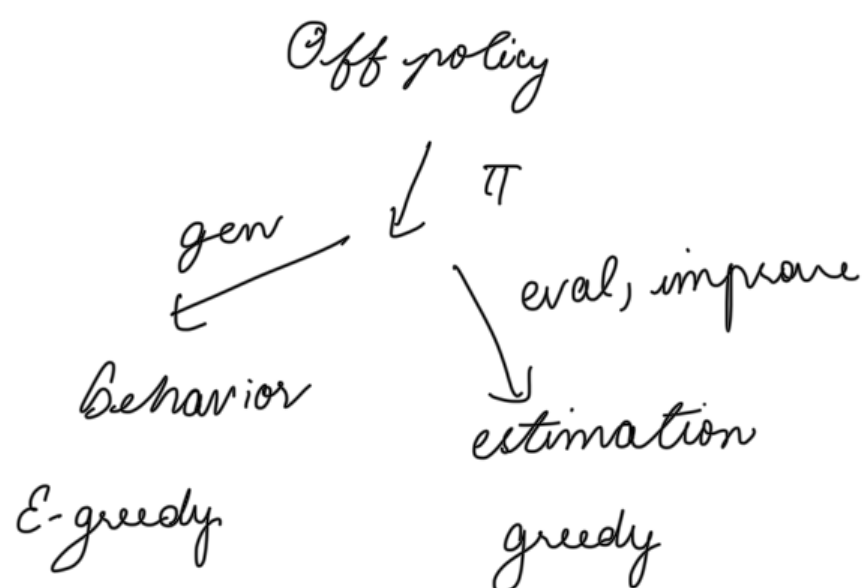
$$G(x) \approx \frac{1}{n} \sum_{i=1}^n x_i$$

$$\begin{aligned} \bar{x}_{n+1} &= \frac{1}{n+1} (x_{n+1} + \bar{x}_n \cdot n) \\ &= \frac{1}{n+1} (x_{n+1} + (n+1)\bar{x}_n - \bar{x}_n) \\ &= \bar{x}_n + \frac{1}{n+1} (x_{n+1} - \bar{x}_n) \\ &= \bar{x}_n + \alpha (x_{n+1} - \bar{x}_n) \end{aligned}$$

$$\text{new est} = \text{oldest} + \alpha (\text{new sample} - \text{old est})$$

TD(0) \rightarrow expectations in Bellman Eq. \rightarrow an average
Q-learning \rightarrow optimality

\rightarrow Q-learning: Off policy TD control



Q-learning Algorithm

initialize $Q(s, a)$ arbitrarily

for each episode

initialize s

for each step of ep

...

can be uniform random or anything you like

$u = \pi(s)$ (Policy derived from Q)
eg ϵ -greedy

estimation policy is always greedy.

$$Q(s,a) \leftarrow \alpha \left[r + \gamma \max_{a'} Q(s',a') - Q(s,a) \right]$$

$s \leftarrow s'$


until s is terminal

→ if ϵ -greedy

only $\epsilon - \frac{\epsilon}{|A(s)|}$ times they differ

SARSA acc to exploratory action
 Q greedy

$\delta = -1$



say 10% times I decide to go to cliff, I fall and update my previous state in SARSA

Q-learning: still update acc to greedy policy.
assumes you will behave greedily in future

Sarsa: ϵ greedy execution in future.
(exploratory)
incorporate cost of exploration

Off Policy Learning

- Target $\pi(a/s) \xrightarrow{\text{eval}} V_{\pi}(s), q_{\pi}(s,a)$

policy

- Behavioural policy

$\mu(a|s)$

eg. if π deterministic

observing somebody else (agent) but want to learn about π

- assumption of coverage:

$$- \pi(a|s) > 0 \rightarrow \mu(a|s) > 0$$

- eg.

π -greedy

μ - ϵ -greedy

→ Importance Sampling:

$$E_{X \sim P} [f(x)] = \sum P(x) f(x)$$

$$= \sum Q(x) \frac{P(x)}{Q(x)} f(x)$$

$$= E_{X \sim Q} \left[\underbrace{\frac{P(x)}{Q(x)}}_{w(x)} f(x) \right]$$

$X \sim P$
but can only
sample acc to
 Q

if I see it often in P but not
 Q it is importance to
give it more weight
and if $P(x) \neq 0 \text{ imp } Q(x) \neq 0$

→ Importance Sampling Ratio

$$\rho_t^T = \frac{\prod_{k=t}^{T-1} \pi\left(\frac{A_k}{S_k}\right) P\left(\frac{S_{k+1}}{S_k, A_k}\right)}{\prod_{k=t}^{T-1} \mu\left(\frac{A_k}{S_k}\right) P\left(\frac{S_{k+1}}{S_k, A_k}\right)}$$

$$= \frac{\prod_{k=t}^{T-1} \pi(A_k | S_k)}{\prod_{k=t}^{T-1} \mu(A_k | S_k)} = \log \sum \frac{\pi}{\mu}$$



weighted average return for V

$$V(s) = \sum_{t=\tau(s)} p_t^T(t) G(t)$$

$$\sum_{t \in \tau(s)} p_t^T(t)$$

Random Variable: $G(t)$

trying to sample & average

$$\frac{N}{D} \quad G(t) = r_t + \gamma r_{t+1} + \dots + \gamma^{T-t-1} r_{T-1}$$

$$S_t \quad A_t \quad S_{t+1} \quad \dots$$